

파일처리 과제: LRU-K 구현

김현준 (2012003954 컴퓨터전공, 한양대학교)

구현내용:

PF_BufPageDesc 및 PF_BufferMgr::InitPageDesc, PF_BufferMgr::GetPage, PF_BufferMgr::UnpinPage 을 패치하여 기존 LRU Page Replacement Algorithm을 LRU-K로 변경했습니다.

LRU-K의 상수 "K"는 src/pf_buffermgr.h 의 K_FOR_LRU_K 정의를 수정하여 설정이 가능합니다.

코드 수정사항:

(GitHub에서 코드 변경사항을 확인하기 위해 생성한 Pull Request(<https://github.com/yoloseem/edubase/pull/1>) 의 commit logs (<https://github.com/yoloseem/edubase/pull/1/commits>) 또는 files changed (<https://github.com/yoloseem/edubase/pull/1/files>) 페이지를 통해 확인하실 수도 있습니다.)

src/pf_buffermgr.h

7 ■■■■ src/pf_buffermgr.h		View	
@@ -15,6 +15,9 @@			
15	15	#ifndef PF_BUFFERMGR_H	
16	16	#define PF_BUFFERMGR_H	
17	17		
	18	++include <time.h>	
	19	++include <deque>	
	20	+	
18	21	#include "pf_internal.h"	
19	22	#include "pf_hashtable.h"	
20	23		
@@ -28,6 +31,9 @@			
28	31	// next.	
29	32	#define INVALID_SLOT (-1)	
30	33		
	34	++ K_FOR_LRU_K is configurable value for LRU-K algorithm	
	35	++define K_FOR_LRU_K 2	
	36	+	
31	37	//	
32	38	// PF_BufPageDesc - struct containing data about a page in the buffer	
33	39	//	
@@ -37,6 +43,7 @@ struct PF_BufPageDesc {			
37	43	int prev;	// prev in the linked list of buffer pages
38	44	int bDirty;	// TRUE if page is dirty
39	45	short int pinCount;	// pin count
	46	+ std::deque<time_t> hists;	// access time histories
40	47	PageNum pageNum;	// page number for this page
41	48	int fd;	// OS file descriptor of this page
42	49	};	

src/pf_buffermgr.cc

24 src/pf_buffermgr.cc

View

@@ -236,6 +236,10 @@ RC PF_BufferMgr::GetPage(int fd, PageNum pageNum, char **ppBuffer,

```
236 236
237 237 // Page is already in memory, just increment pin count
238 238 bufTable[slot].pinCount++;
239 + bufTable[slot].hists.push_back(time(NULL));
240 + if (bufTable[slot].hists.size() > K_FOR_LRU_K)
241 + // Keep only K(in LRU-K) histories
242 + bufTable[slot].hists.pop_front();
```

```
239 243 #ifdef PF_LOG
240 244 sprintf (psMessage, "Page found in buffer. %d pin count.\n",
241 245 bufTable[slot].pinCount);
```

@@ -385,6 +389,7 @@ RC PF_BufferMgr::UnpinPage(int fd, PageNum pageNum)

```
385 389 #endif
386 390
387 391 // If unpinning the last pin, make it the most recently used page
392 + bufTable[slot].hists.pop_front();
388 393 if (--(bufTable[slot].pinCount) == 0) {
389 394 if ((rc = Unlink(slot)) ||
390 395 (rc = LinkHead (slot)))
```

@@ -784,10 +789,19 @@ RC PF_BufferMgr::InternalAlloc(int &slot)

```
784 789 }
785 790 else {
786 791
787 - // Choose the least-recently used page that is unpinned
788 - for (slot = last; slot != INVALID_SLOT; slot = bufTable[slot].prev) {
789 - if (bufTable[slot].pinCount == 0)
790 - break;
792 + // Choose the least-recently used page that is unpinned
793 + int &candidate = slot, accessedTime = -1;
794 + for (candidate = last;
795 + candidate != INVALID_SLOT;
796 + candidate = bufTable[candidate].prev) {
797 + if (bufTable[slot].pinCount == 0) {
798 + if (accessedTime == -1 ||
799 + bufTable[slot].hists.front() < accessedTime) {
800 + slot = candidate;
801 + accessedTime = bufTable[candidate].hists.back();
802 + break;
803 + }
804 + }
```

```
791 805 }
792 806
793 807 // Return error if all buffers were pinned
```

@@ -910,6 +924,8 @@ RC PF_BufferMgr::InitPageDesc(int fd, PageNum pageNum, int slot)

```
910 924 bufTable[slot].pageNum = pageNum;
911 925 bufTable[slot].bDirty = FALSE;
912 926 bufTable[slot].pinCount = 1;
927 + bufTable[slot].hists = deque<time_t>();
928 + bufTable[slot].hists.push_back(time(NULL));
```

```
913 929
914 930 // Return ok
915 931 return (0);
```

테스트 결과:

pf_test1

```
1. red008@proxy: ~/edubase/src (ssh)
Got page: 57 57
Got page: 58 58
Got page: 59 59
Got page: 60 60
Got page: 61 61
Got page: 62 62
Got page: 63 63
Got page: 64 64
Got page: 65 65
Got page: 66 66
Got page: 67 67
Got page: 68 68
Got page: 69 69
Got page: 70 70
Got page: 71 71
Got page: 72 72
Got page: 73 73
Got page: 74 74
Got page: 75 75
Got page: 76 76
Got page: 77 77
Got page: 78 78
Got page: 79 79
EOF reached
Closing and destroying both files
PF Layer Statistics
-----
Total number of calls to GetPage Routine: 702
  Number found: 23
  Number not found: 679
-----
Number of read requests: 679

Number of write requests: 339
-----
Number of flushes: 16
-----
Verifying the statistics for buffer manager: Correct!
Testing hash table. Inserting entries
Searching for entries
Deleting entries in reverse order
Ensuring all entries were deleted
Ending PF layer test.

red008@proxy:~/edubase/src$
```

pf_test2

```
1. red008@proxy: ~/edubase/src (ssh)

Number of write requests: 339
-----
Number of flushes: 16
-----
Verifying the statistics for buffer manager: Correct!
Testing hash table. Inserting entries
Searching for entries
Deleting entries in reverse order
Ensuring all entries were deleted
Ending PF layer test.

red008@proxy:~/edubase/src$ ./pf_test2
*****
Starting PF layer test.
-----
Creating file: file1
Opening file: file1
Allocating 40 pages.
Asking for the same pages again.
Verifying the statistics for buffer manager: Correct!
Unpinning pages.
Verifying the write statistics for buffer manager: Correct!
Allocating an additional 40 pages to clearout the buffer pool.
Now asking for the original pages again.
Verifying that pages were not found in buffer pool: Correct!
Flushing the File handle to disk.
Flushing the File handle to disk. (Again)
Testing number of pages written to disk: Correct!
PF Layer Statistics
-----
Total number of calls to GetPage Routine: 80
    Number found: 40
    Number not found: 40
-----
Number of read requests: 40

Number of write requests: 80
-----
Number of flushes: 4
-----
Ending PF layer test.
*****

red008@proxy:~/edubase/src$
```

pf_test3

```
1. red008@proxy: ~/edubase/src (ssh)

-----
Creating file: file1
Opening file: file1
Allocating 40 pages.
Asking for the same pages again.
Verifying the statistics for buffer manager: Correct!
Unpinning pages.
Verifying the write statistics for buffer manager: Correct!
Allocating an additional 40 pages to clearout the buffer pool.
Now asking for the original pages again.
Verifying that pages were not found in buffer pool: Correct!
Flushing the File handle to disk.
Flushing the File handle to disk. (Again)
Testing number of pages written to disk: Correct!
PF Layer Statistics
-----
Total number of calls to GetPage Routine: 80
  Number found: 40
  Number not found: 40
-----
Number of read requests: 40

Number of write requests: 80
-----
Number of flushes: 4
-----
Ending PF layer test.
*****

red008@proxy:~/edubase/src$ ./pf_test3
Starting PF Chunk layer test.
Note: Statistics are turned on.
Asking for 10 chunks from the buffer manager: Pass
Verifying the contents of 10 chunks from the buffer manager: Pass
Disposing of 5 chunks from the buffer manager: Pass
Verifying the contents of 10 chunks from the buffer manager: Pass
Asking for 35 chunks from the buffer manager: Pass
Asking for 1 chunks from the buffer manager: Pass
Disposing of 1 chunks from the buffer manager: Pass
Disposing of 35 chunks from the buffer manager: Pass
Asking for 25 chunks from the buffer manager: Pass
Verifying the contents of 25 chunks from the buffer manager: Pass
Ending PF Chunk layer test.

red008@proxy:~/edubase/src$
```